

# Minor Thesis

## Development and Analysis of Barrier Protocols

Ronny Brendel (<http://automaton2000.com>)  
Responsible Professor: Prof. Dr. Christel Baier  
Supervisor: Dr. Sascha Klüppelholz

Dresden, 2014-01-08

- Introduction

Basics, Motivation

- Protocols

Central Counter, B1 Barrier

- Modelling

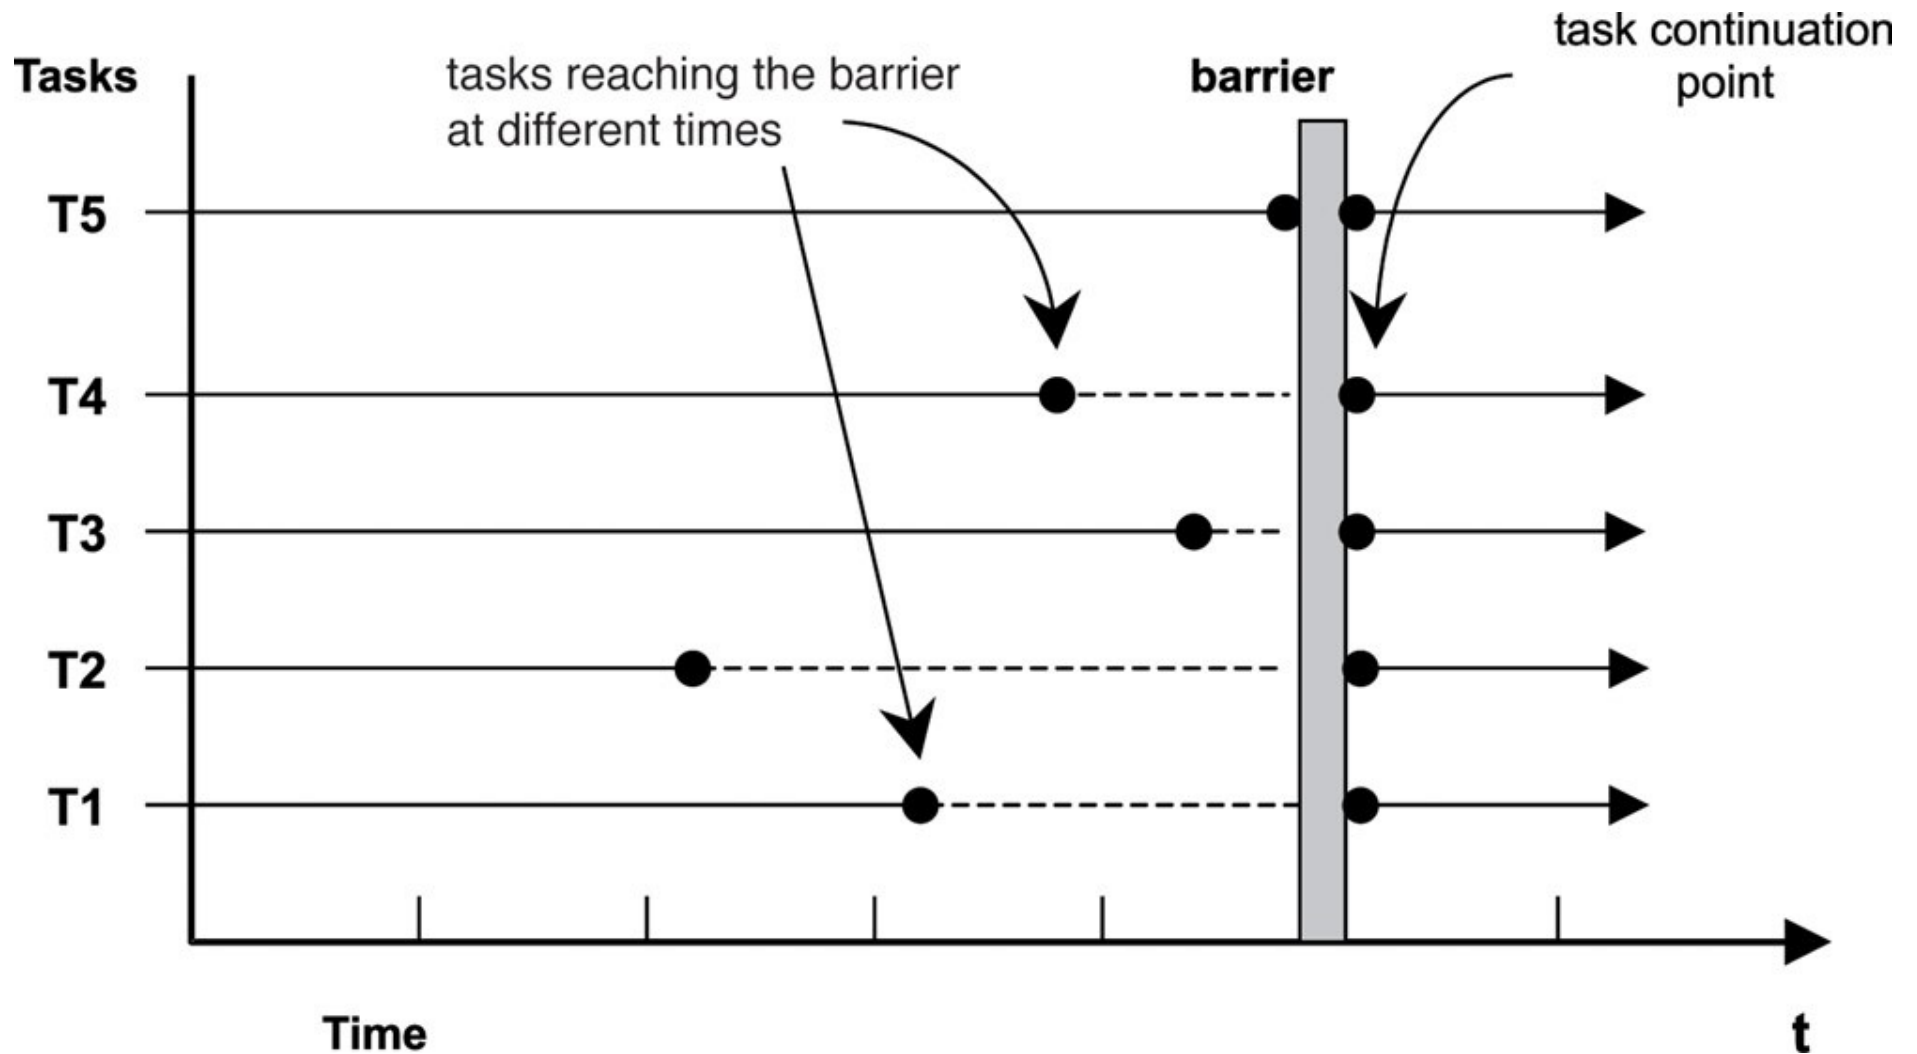
Shared Variable, Protocols

- Analysis


Functional, Quantitative

- Conclusion, Future Work, Sources

# Basics ◁ Introduction

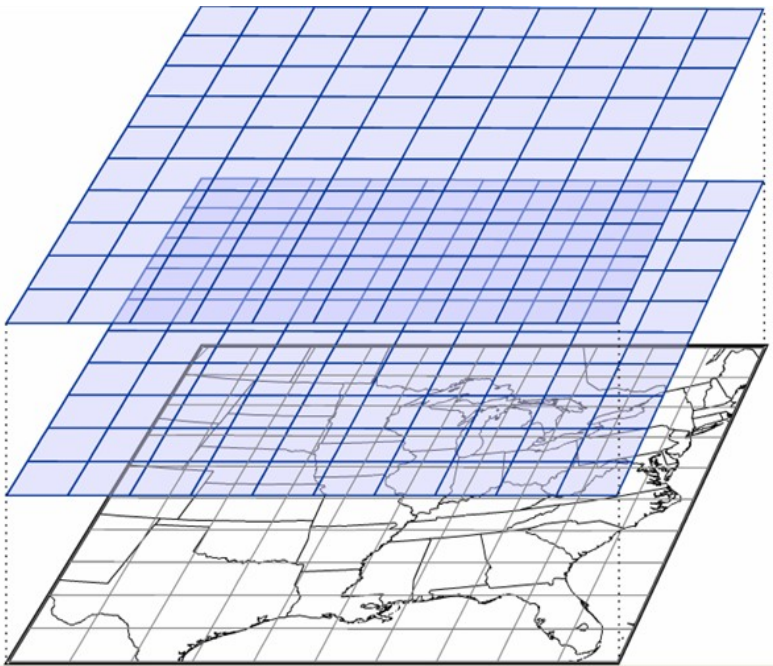


# Basics ◁ Introduction

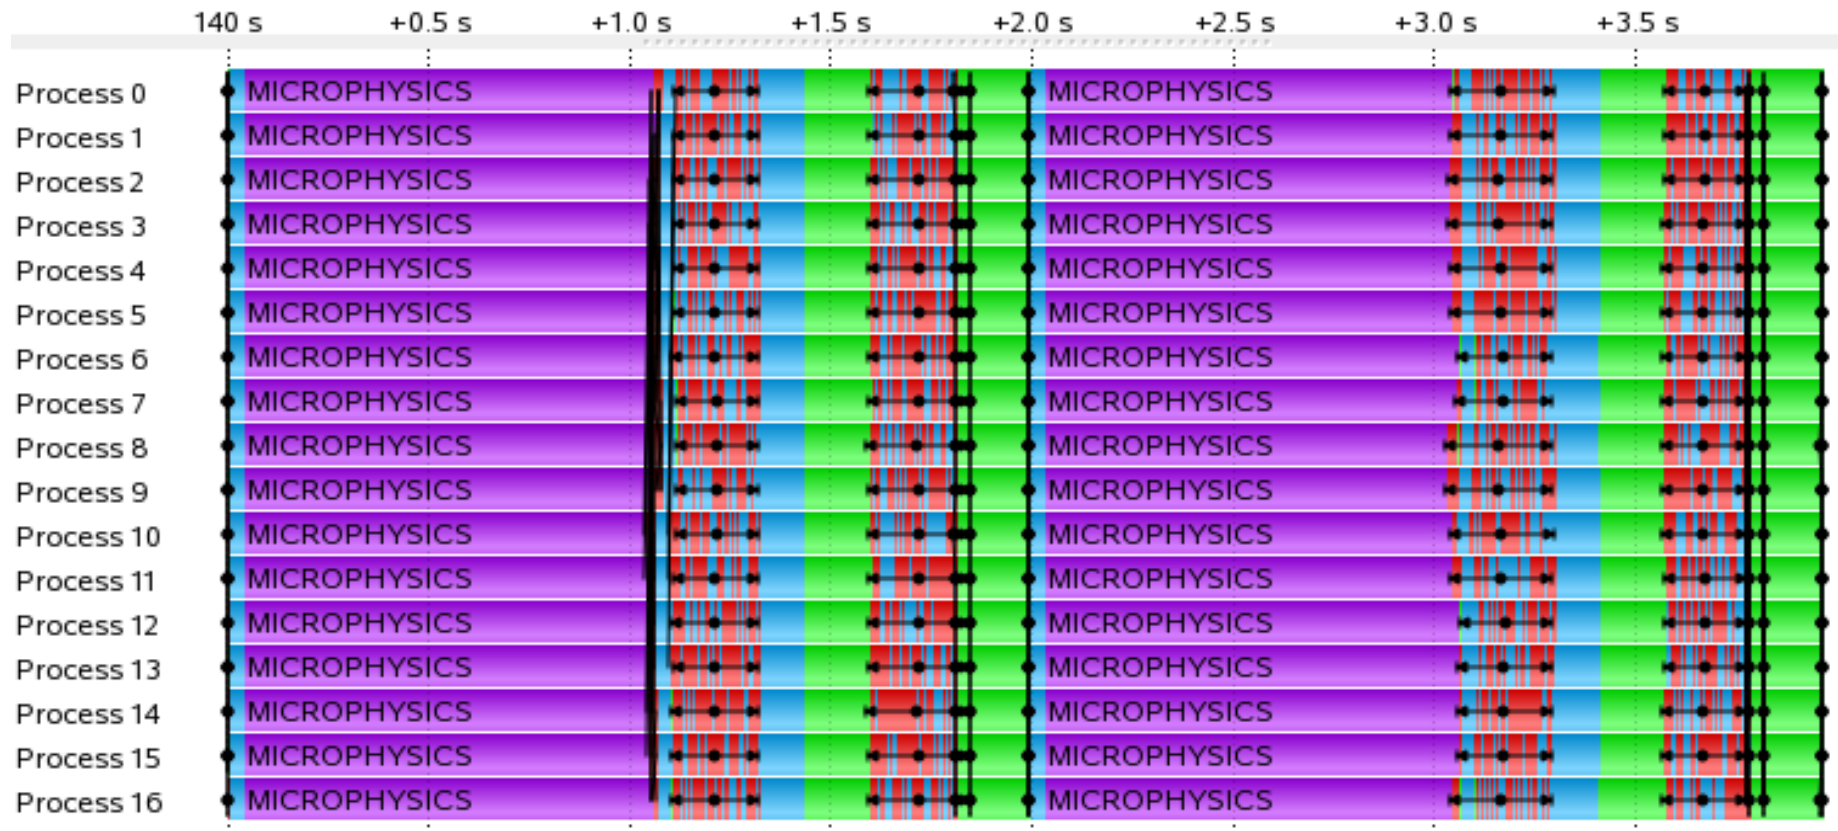


# Basics $\triangleleft$ Introduction

---



# Basics ◁ Introduction

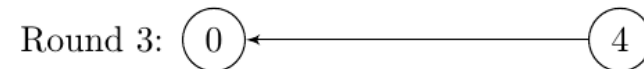
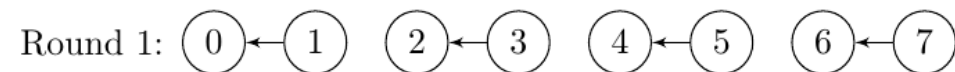


# Basics ◁ Introduction


---

- Usual Implementations include
  - Central Counter Barrier (atomic increment)
  - Hierarchical approaches

**Gather:**



**Broadcast:**



see sources [3]

# Motivation $\triangleleft$ Introduction

---

- Today's Barrier Protocols have been invented long ago
- Probabilistic Write/Copy-Select (pW/CS)
  - Concurrent protocols are unnecessarily strict
  - Relieving strictness can improve performance
  - Complexity of modern computers makes the timing of concurrent interaction effectively random. Employ the tools of probability theory for designing/analysing protocols



# Motivation $\triangleleft$ Introduction

---

- Tests/Benchmarks
  - not exhaustive
  - not arbitrarily fine-grained, incurs overhead
  - testing probabilistic algorithms is hard
- Model checking
  - exhaustive
  - arbitrarily fine-grained, no overhead

# Motivation ◁ Introduction

---

- Improve barrier protocols
- Using the principles behind pW/CS lock
- Analysis through model checking

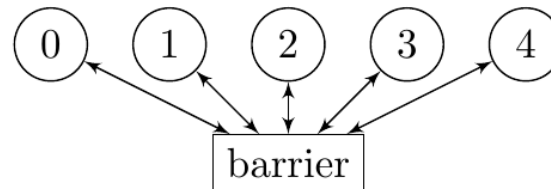
# Central Counter Bar. $\triangleleft$ Protocols

```
shared variables: integer barrier := threadCount
```

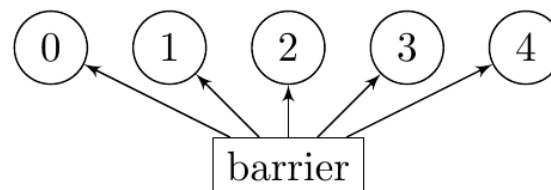
```
atomic{barrier := barrier - 1}
```

```
wait until barrier = 0
```

Atomic decrement:



Repeated reading:



# B1 Barrier $\triangleleft$ Protocols

---

```
shared variables: boolean barrier[threadCount]
local variables:  integer i
initialisation:   barrier[*] := false
```

---

```
barrier[threadIndex] := true
```

```
i := 0
while i < threadCount {
    if barrier[i] = false {
        i := -1
    }
    i = i + 1
}
```

# Modelling

---

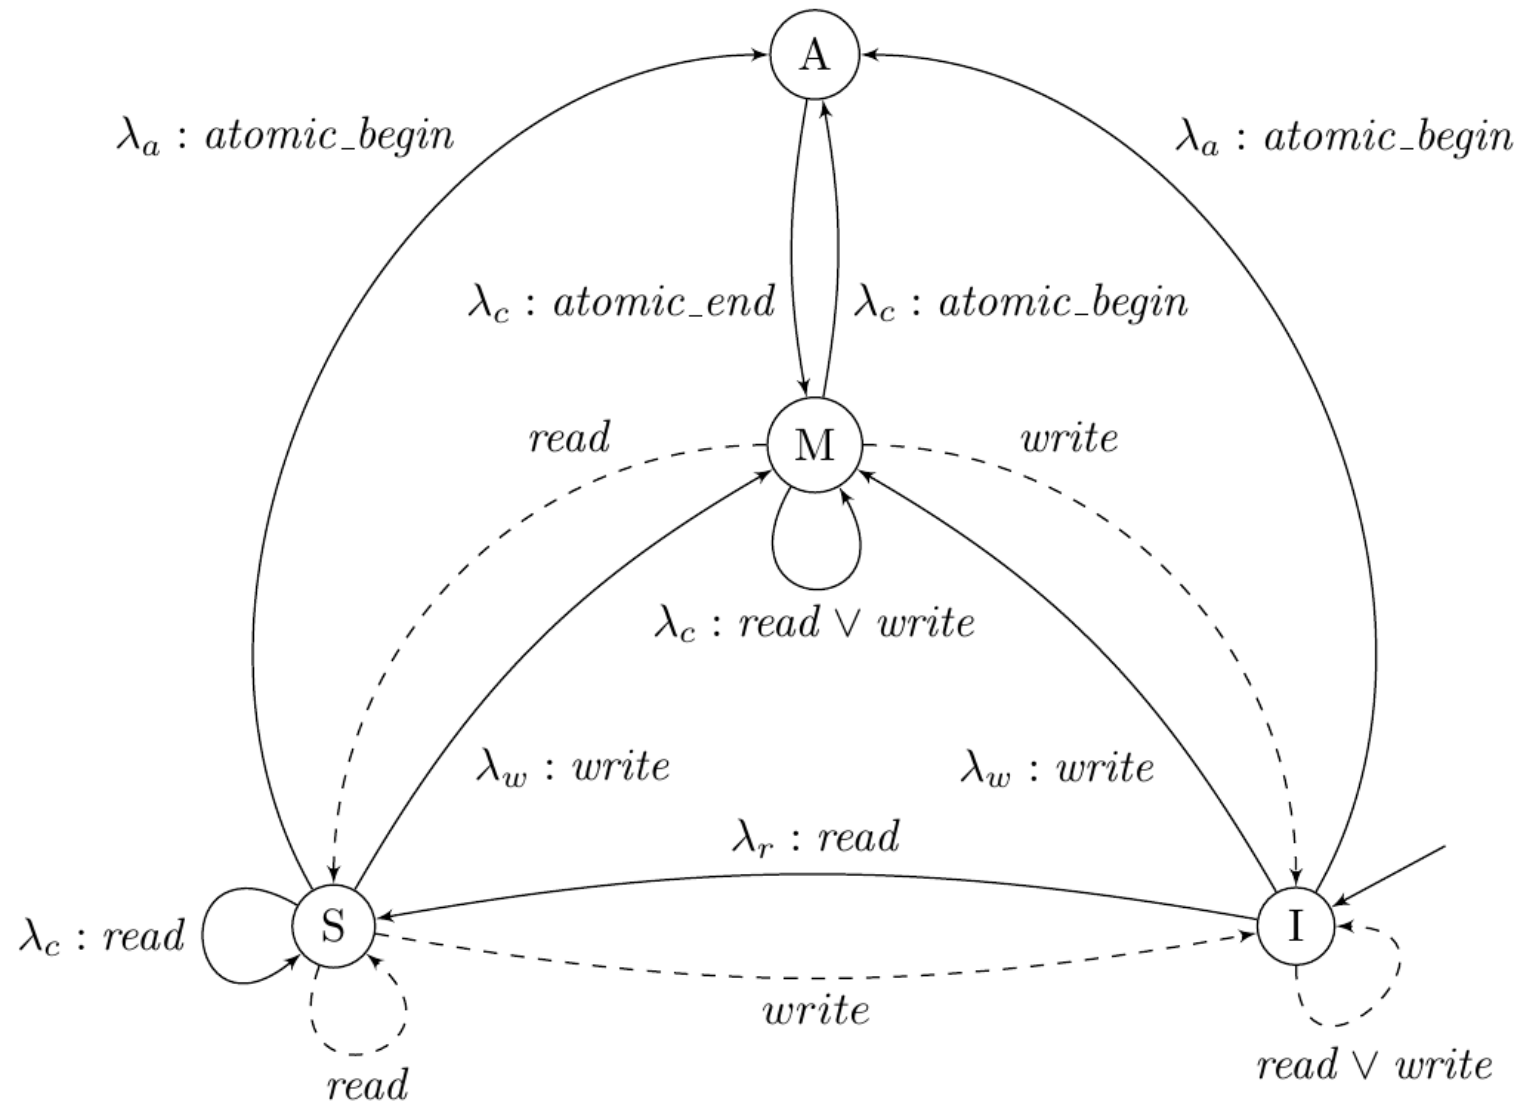
- Functional
  - Non-deterministic transition system + LTL
  - SPIN
  - detailed model to reveal all possible mistakes
- Quantitative
  - CTMC + CSL/CSRL
  - PRISM
  - reduced to just costly/important transitions, no reinitialisation

# Shared Variable ◁ Modelling

---

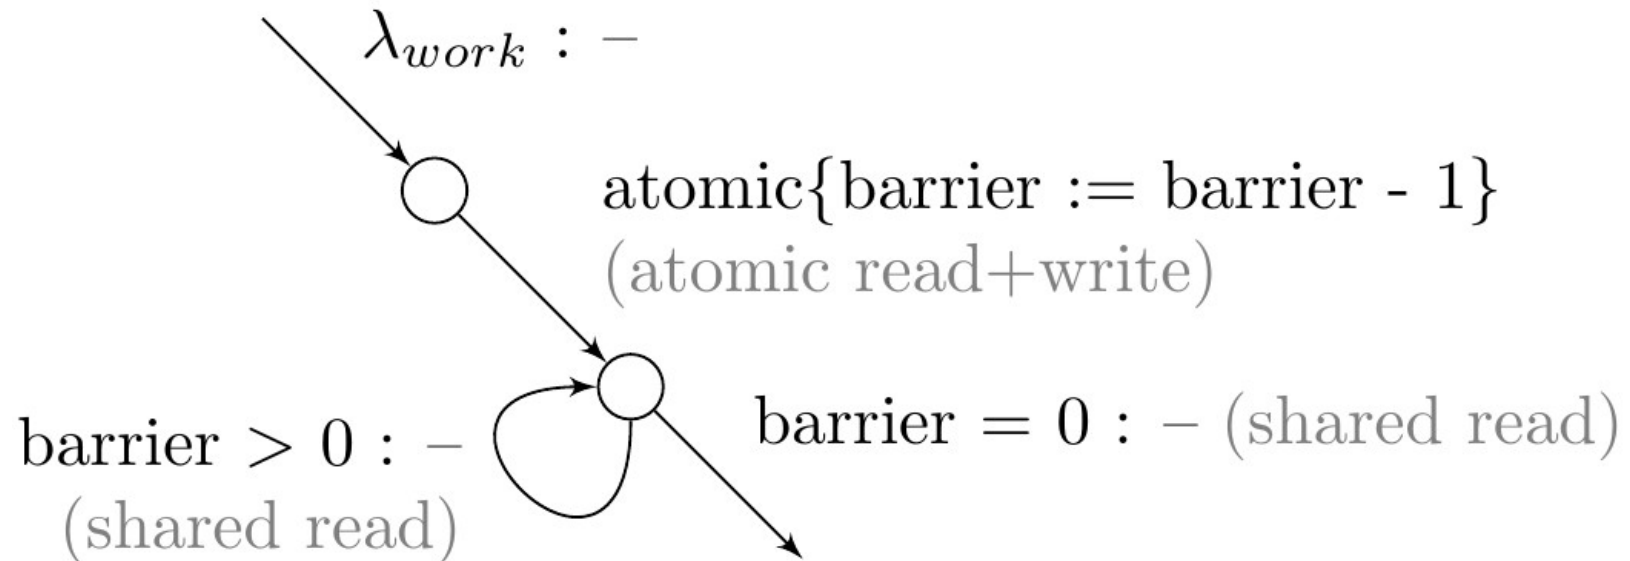
- Synchronisation is about exchanging information, i.e. sharing memory
- Very small information → Timing dominated by memory access latency
- Memory access is cached → We have to model caching
- We identify a shared variable with the cache line it resides on
- MSI protocol + atomic operations

# Shared Variable $\triangleleft$ Modelling



# Central Counter Bar. $\triangleleft$ Modelling

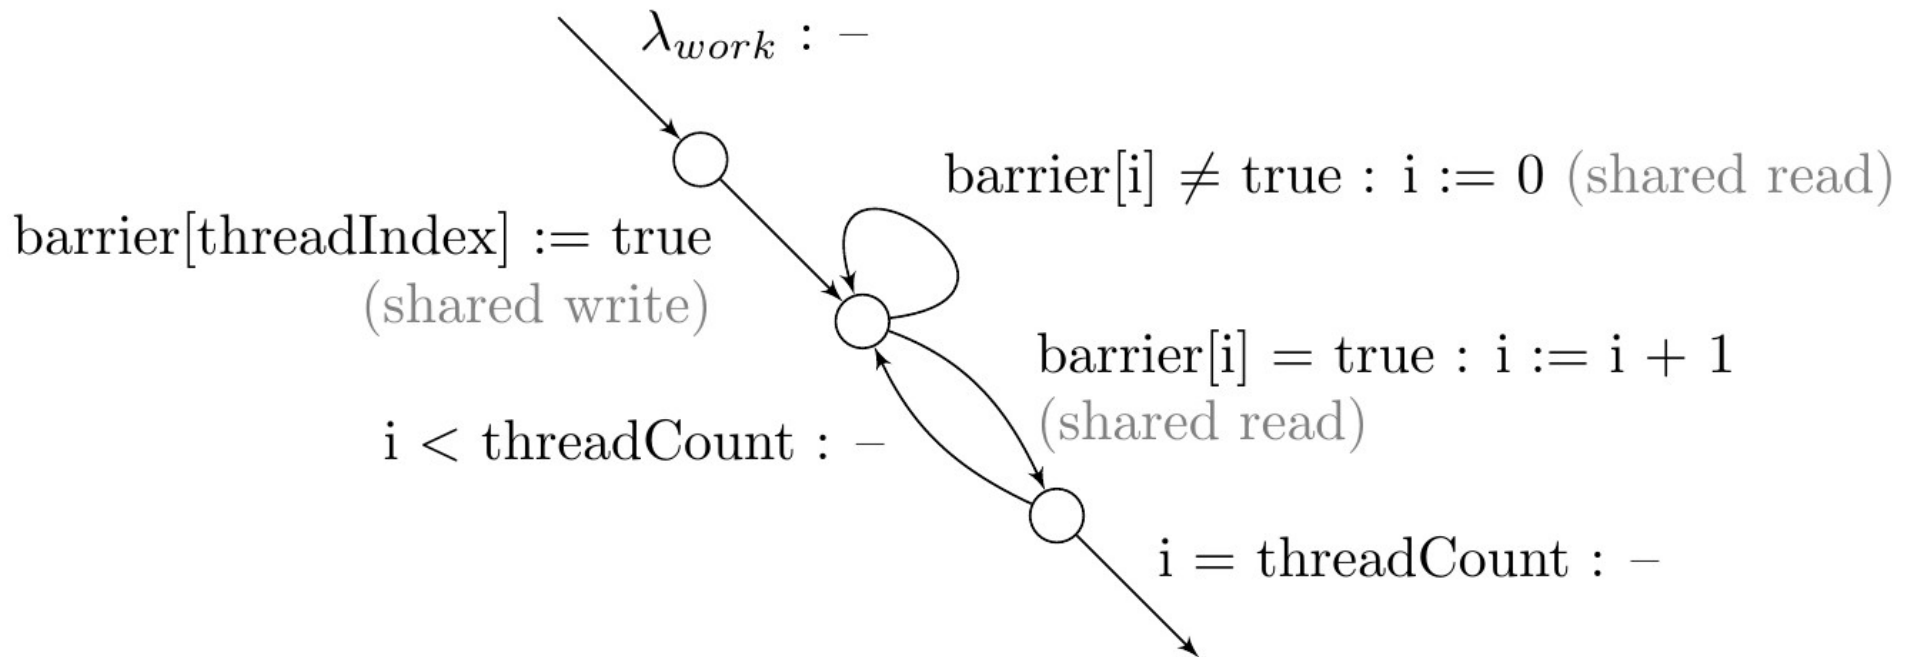
---





# B1 Barrier $\triangleleft$ Modelling

---



# Functional $\triangleleft$ Analysis

---

- “A thread may only exit the barrier if all threads have entered it”

$$\Box(one\_left \implies all\_entered)$$

- “If all threads entered the barrier, each one leaves it in a finite amount of time”

$$\Box(all\_entered \implies \Diamond all\_left)$$

# Functional $\triangleleft$ Analysis


---

- “A thread may only exit the barrier if all threads have entered it”

$$\square(one\_left \implies all\_entered)$$

- “If all threads entered the barrier, each one leaves it in a finite amount of time”

$$\square(all\_entered \implies \diamond all\_left)$$



# Quantitative $\triangleleft$ Analysis

---

- Interesting values
  - time in cycles (2.5GHz clock speed)
  - energy consumption in joule / watt
  - implemented using CTMC time and rewards
- Parameters to variate
  - Number of threads
  - Rate of the initial (work) distribution

# Quantitative $\triangleleft$ Analysis

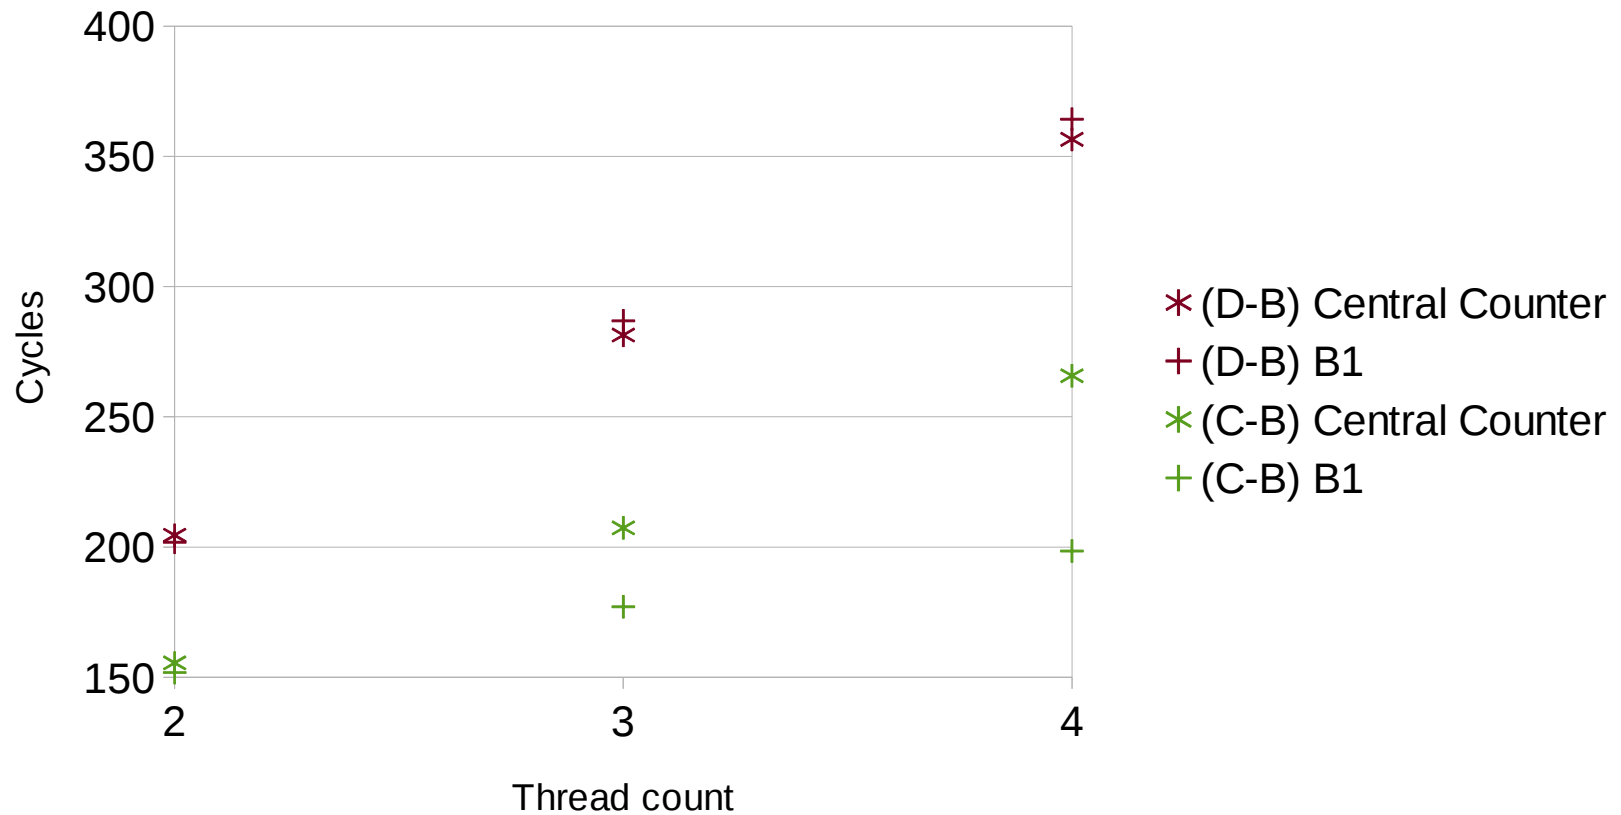
---

- Points in time when to measure
  - (A) First thread entered the barrier
  - (B) Last thread entered
  - (C) First thread left
  - (D) Last thread left
  - (W) Last thread finished writing
- We take the reachability reward wrt the proper reward functions at these points in time

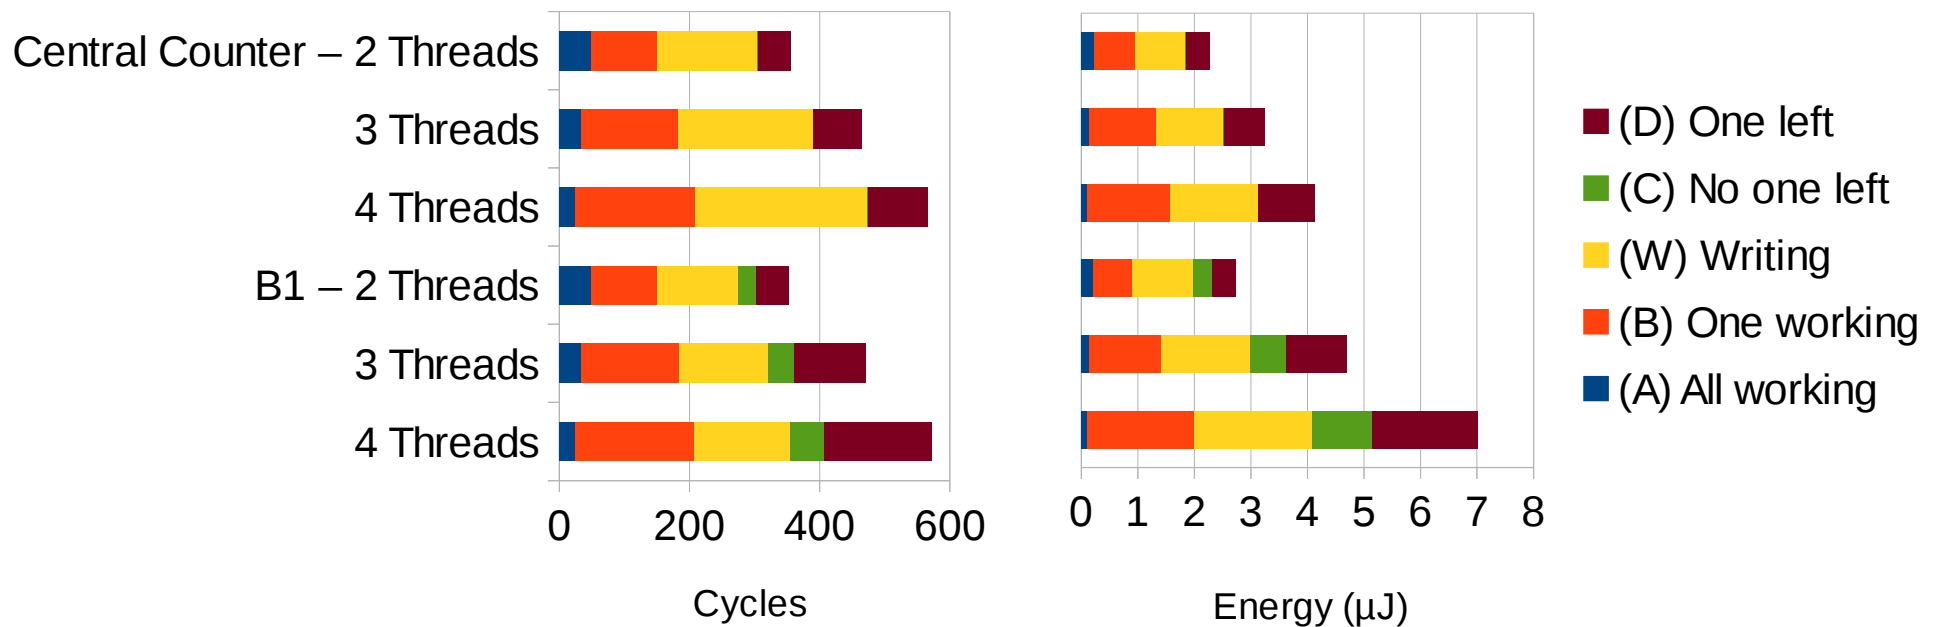
# Quantitative $\triangleleft$ Analysis

---

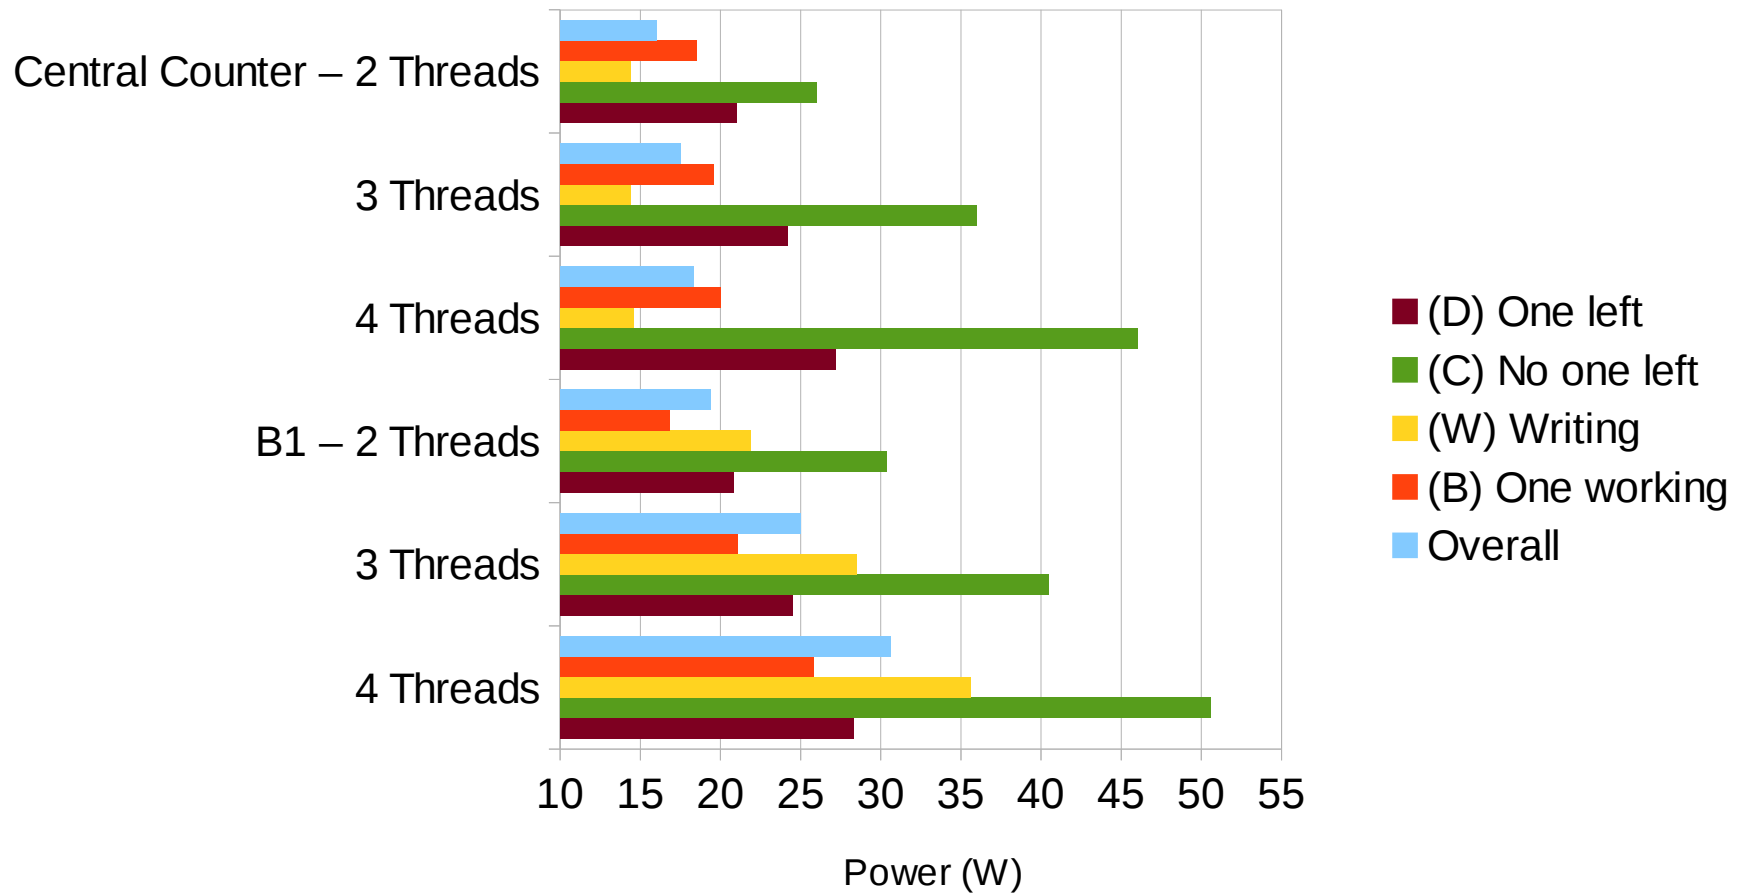
- small work period (100 cycles)



# Quantitative $\triangleleft$ Analysis



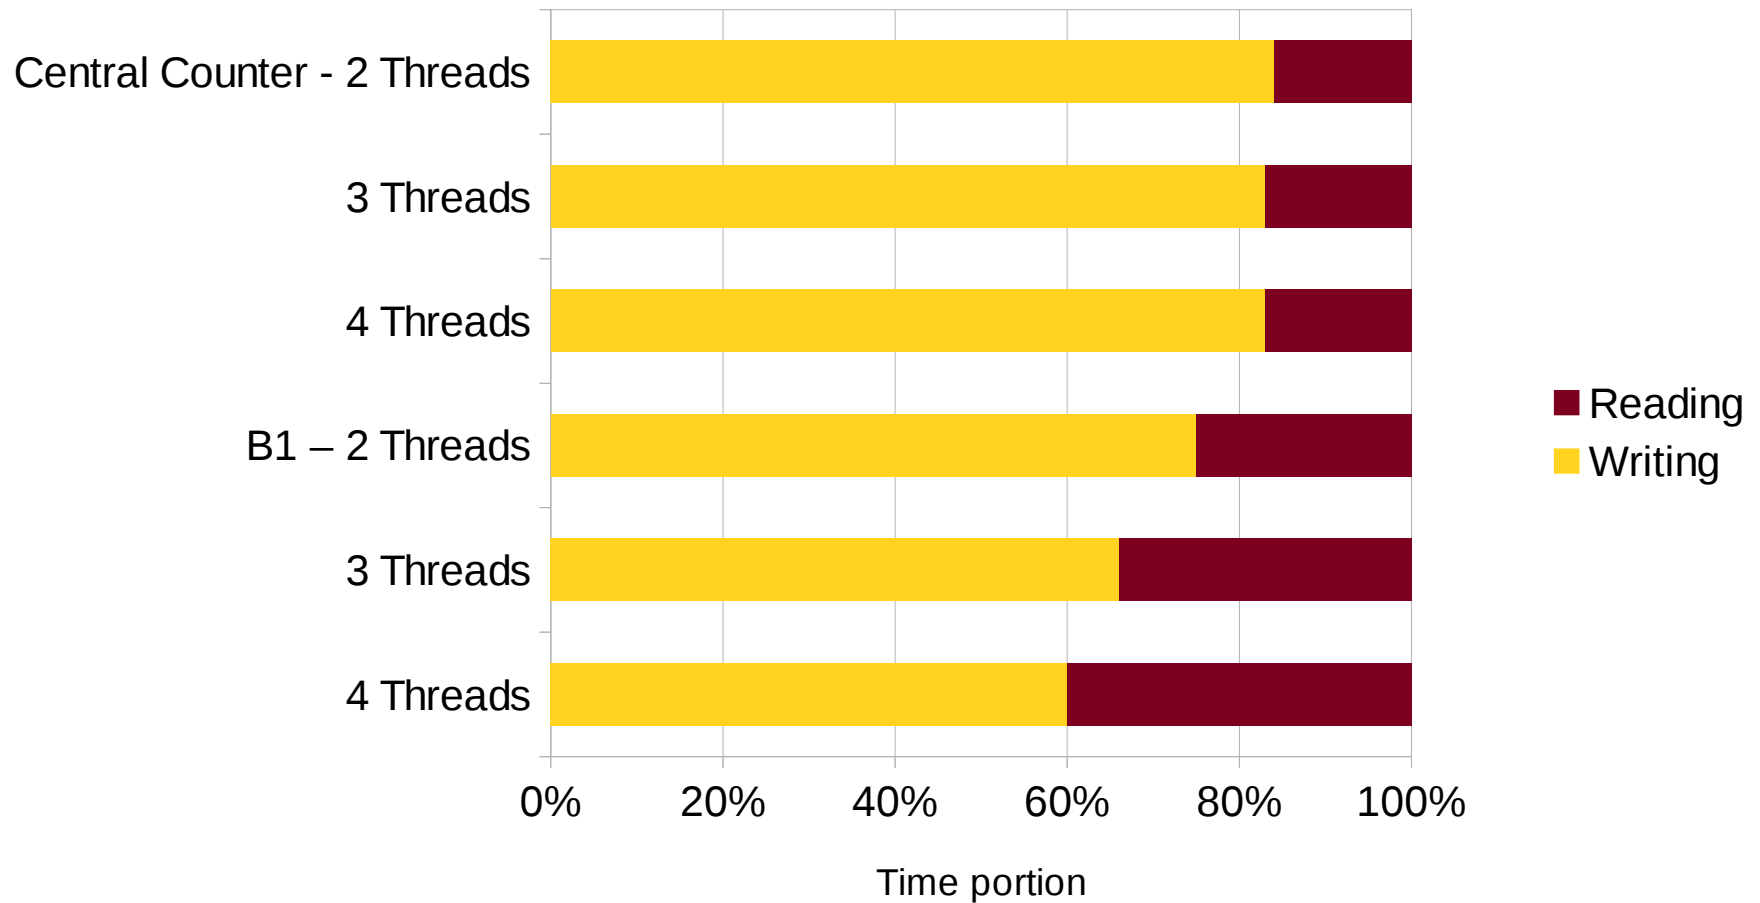
# Quantitative $\triangleleft$ Analysis





# Quantitative Analysis

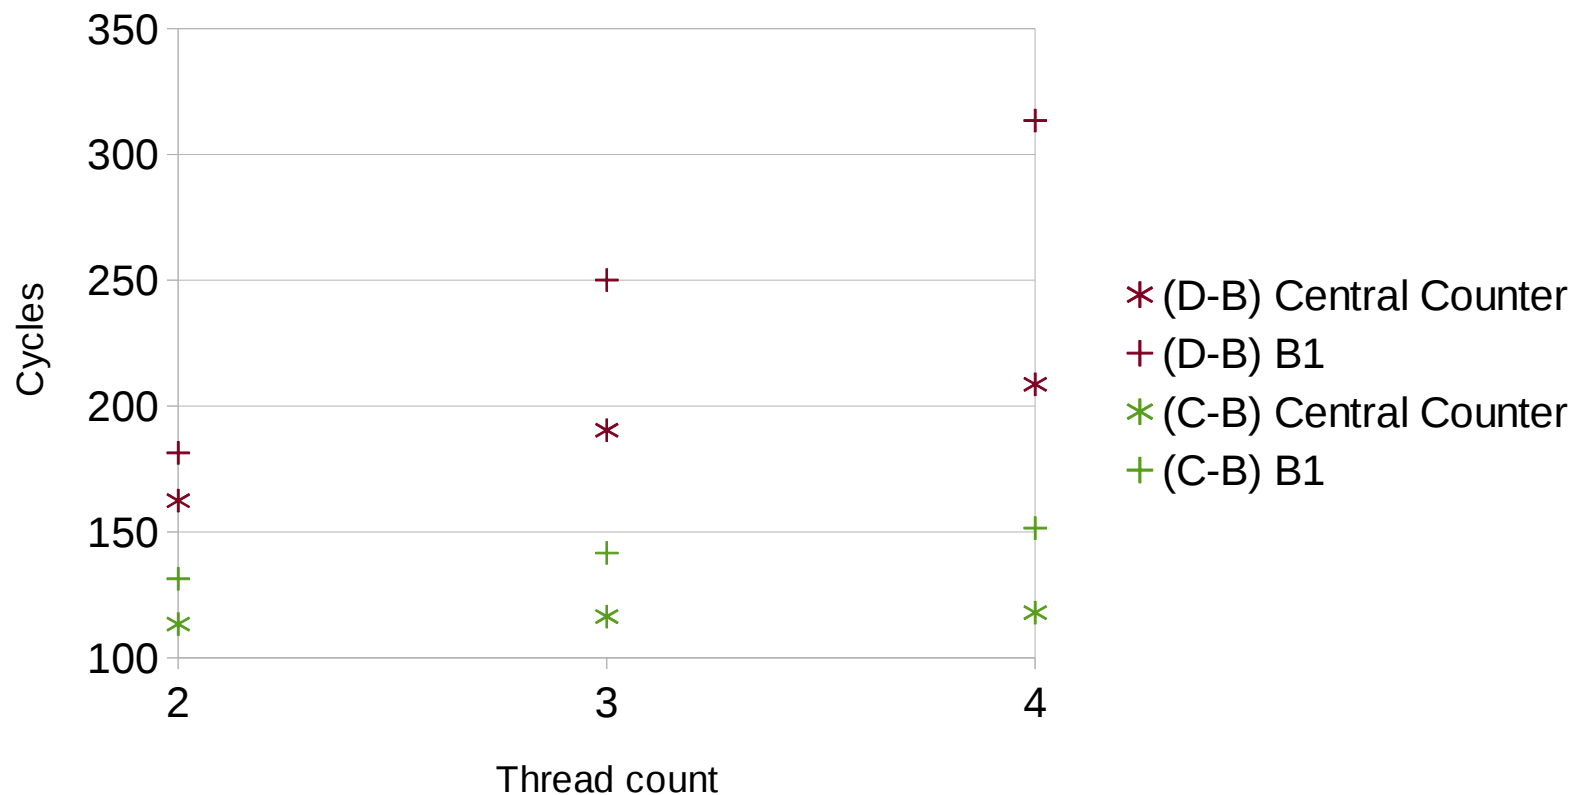
---



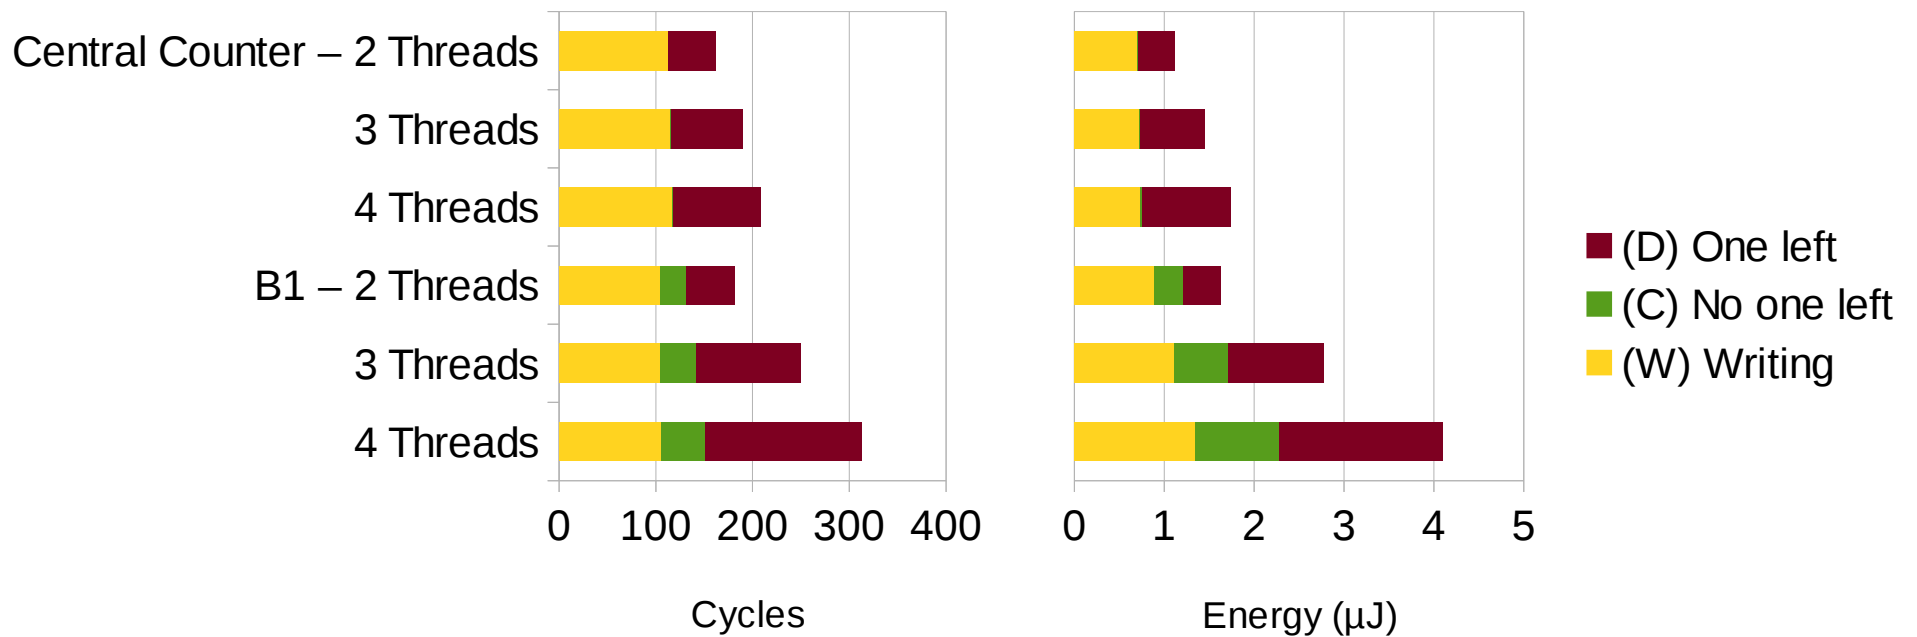
# Quantitative ◁ Analysis

---

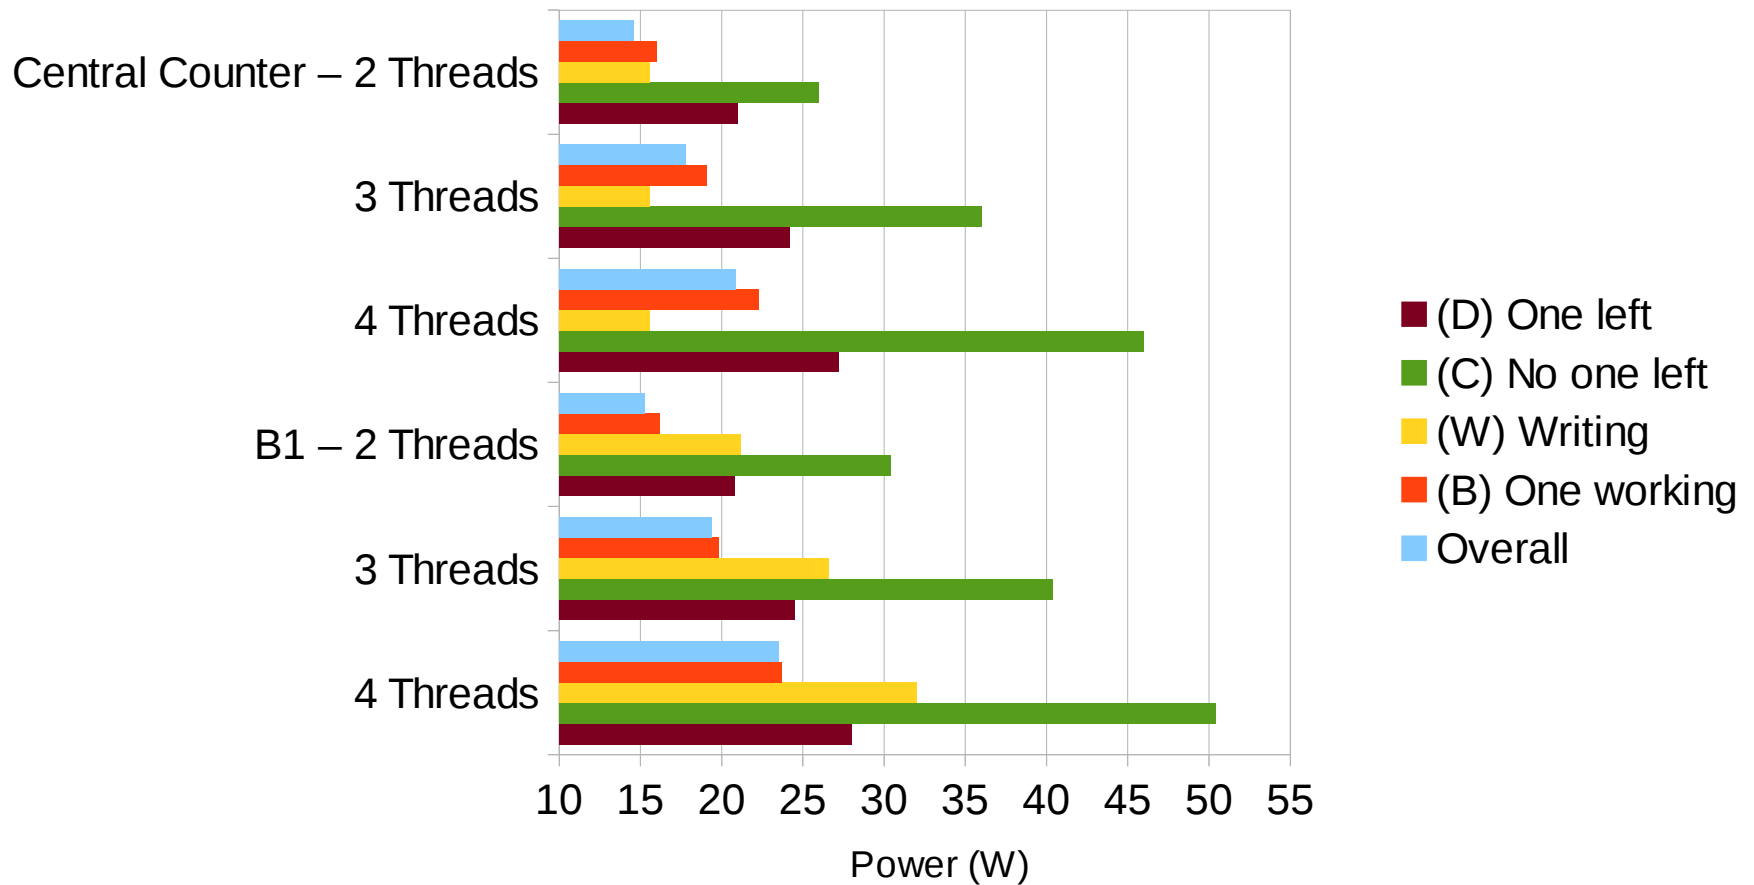
- large work period (1000 cycles)



# Quantitative $\triangleleft$ Analysis



# Quantitative $\triangleleft$ Analysis



# Conclusion

---

- Introduced innovative barrier protocols
  - + No atomic operations or locks required
  - + Competitive performance
- Principles of pW/CS locks apt to improve synchronisation performance
- Quantitative model checking enables exhaustive, fine-grained analysis beyond the capability of tests and benchmarks

# Future Work

---

- Analyse protocols using measurement
- Invent more protocols
  - Variations of existing
  - Remote write-based
- Extend model checking
  - More processes/threads
  - More detail
    - Cache protocols, cache hierarchies
    - Limited bandwidth and other influences

- [1] Probabilistic write copy select, Paper,  
In 13th Real-Time Linux Workshop, pages 195–206, Oct. 2011
- [2] A probabilistic quantitative analysis of  
probabilistic-write/copy-select, Paper,  
In NASA Formal Methods, pages 307–321. Springer, 2013.
- [3] Evaluation of publicly available Barrier-Algorithms and  
Improvement of the Barrier-Operation for large-scale  
Cluster-Systems with special Attention on InfiniBand Networks  
<http://hlor.inf.ethz.ch/publications/index.php?pub=12>
- [4] PRISM, Website, 13-03-019  
<http://www.prismmodelchecker.org>
- [5] SPIN, Website, 13-01-08  
<http://spinroot.com>

# Not Covered in the Presentation

---

- Survey of means to implement barriers
- Barrier building blocks
- Dissemination Barrier, MGB and B2 Barrier
- more in-depth analysis of the protocols



# Thank you!

Slides and report are available at

<http://automaton2000.com/barrier-slides.pdf>

<http://automaton2000.com/barrier-minor-thesis.pdf>

# Minor Thesis

## Development and Analysis of Barrier Protocols

Ronny Brendel (<http://automaton2000.com>)

Responsible Professor: Prof. Dr. Christel Baier

Supervisor: Dr. Sascha Klüppelholz


Dresden, 2014-01-08

# Content

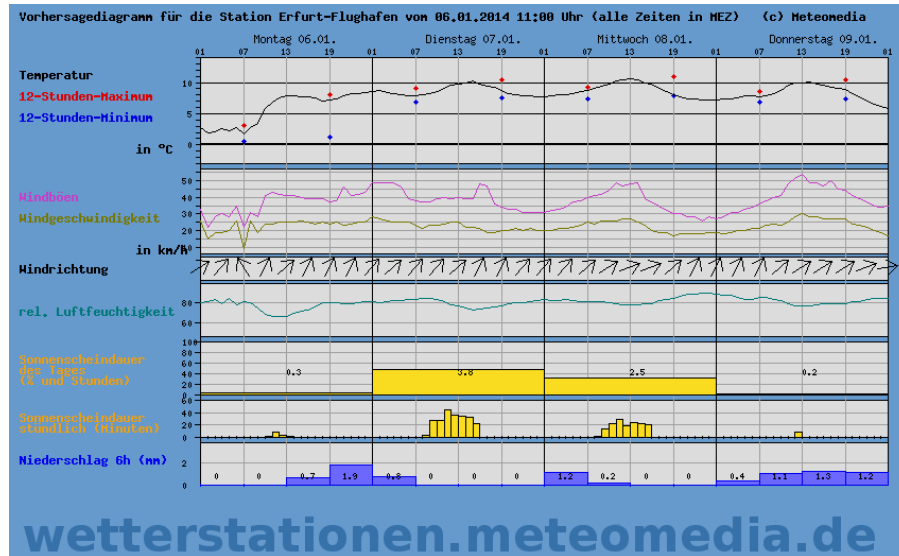
---

- Introduction
  - Basics, Motivation
- Protocols
  - Central Counter, B1 Barrier
- Modelling
  - Shared Variable, Protocols
- Analysis
  - Functional, Quantitative
- Conclusion, Future Work, Sources

# Basics ◁ Introduction

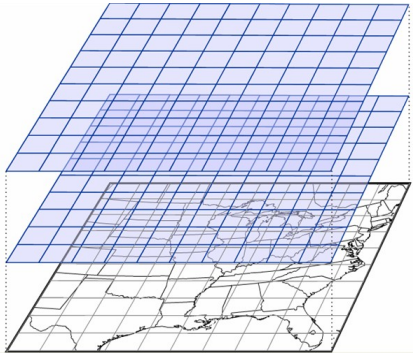


# Basics ◀ Introduction

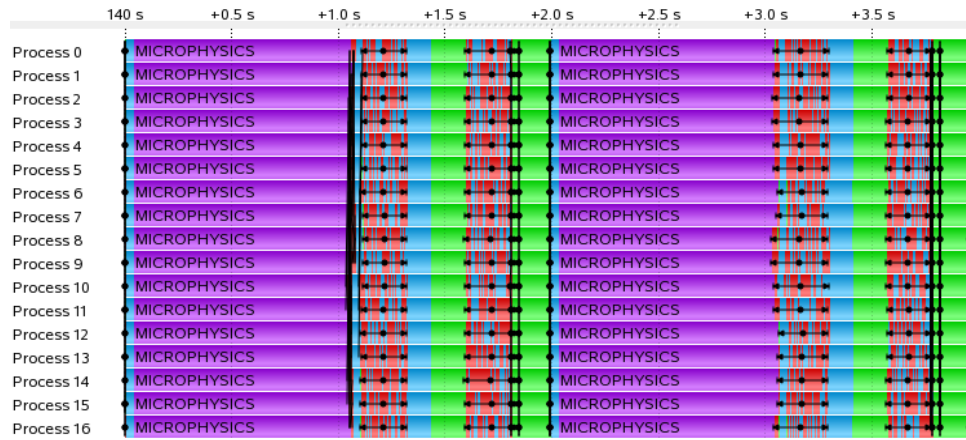


# Basics ◁ Introduction

---



# Basics ◀ Introduction



Lock-step / Gleichschritt

# Basics ◁ Introduction

---

- Usual Implementations include
  - Central Counter Barrier (atomic increment)
  - Hierarchical approaches

**Gather:**

Round 1: (0) ← (1)   (2) ← (3)   (4) ← (5)   (6) ← (7)

Round 2: (0) ← (2)   (4) ← (6)

Round 3: (0) ← (4)

**Broadcast:**

Round 1: (0) → (4)

Round 2: (0) → (2)   (4) → (6)

Round 3: (0) → (1)   (2) → (3)   (4) → (5)   (6) → (7)

see sources [3]



## Motivation ◁ Introduction

---

- Today's Barrier Protocols have been invented long ago
- Probabilistic Write/Copy-Select (pW/CS)
  - Concurrent protocols are unnecessarily strict
  - Relieving strictness can improve performance
  - Complexity of modern computers makes the timing of concurrent interaction effectively random. Employ the tools of probability theory for designing/analysing protocols

see sources [1,2]

8

- Dissemination 1988
- Central Counter noch aelter
- Neue Variationen und andere Technologien, aber im Kern das selbe (RMA Dissemination, n-way Dissemination)
- Neuere workshop paper von Nicholas Mc Guire
  - locking ist kompliziert
  - locking/atomic ops skalieren schlecht
  - pW/CS lock = Alternatives locking Technik
    - Replikation
    - Fehlertoleranz
    - Wahrscheinlichkeit des Fehlers so gering dass im Durchschnitt die performance besser ist.

## Motivation ◁ Introduction

---

- Tests/Benchmarks
  - not exhaustive
  - not arbitrarily fine-grained, incurs overhead
  - testing probabilistic algorithms is hard
- Model checking
  - exhaustive
  - arbitrarily fine-grained, no overhead

9

- Nachteile
  - Aufwand. Messen vs Simulieren/Analysieren
  - Skaliert schlecht
  - Ergebnisqualitaet ist abhaengig von der Modellqualitaet

## Motivation ◁ Introduction

---

- Improve barrier protocols
- Using the principles behind pW/CS lock
- Analysis through model checking

10

- ein Gaengiger Algorithmus vs einen Neuen

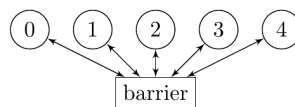
# Central Counter Bar. $\triangleleft$ Protocols

```
shared variables: integer barrier := threadCount
```

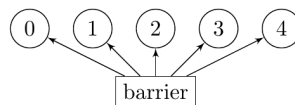
```
atomic{barrier := barrier - 1}
```

```
wait until barrier = 0
```

Atomic decrement:



Repeated reading:



11

!!! Wir behandeln im Vortrag nur shared memory !!!

$\text{ceil}(\log(2, n))$  bytes

## B1 Barrier ◁ Protocols

---

```
shared variables: boolean barrier[threadCount]
local variables: integer i
initialisation:   barrier[*] := false

barrier[threadIndex] := true

i := 0
while i < threadCount {
    if barrier[i] = false {
        i := -1
    }
    i = i + 1
}
```

12

- $n * 64$  bytes - linear
- aehnlich wie central counter vom ablauf her
  - einmal schreiben
  - oft lesen
- schnelleres schreiben, weil keine atomic ops, gleichzeitig moeglich
- langsameres lesen, weil mehrere variablen gelesen werden muessen.
- Da nach dem schreiben einer variable nie wieder in die gleiche Variable geschrieben wird, wird es in jedem thread gecacht und ist dann schnell verfuegbar

# Modelling

---

- Functional
  - Non-deterministic transition system + LTL
  - SPIN
  - detailed model to reveal all possible mistakes
- Quantitative
  - CTMC + CSL/CSRL
  - PRISM
  - reduced to just costly/important transitions, no reinitialisation

see sources [4,5]

13

- Folgend hauptsaechlich quantitatives Modelling
- Funktionales ist aehnlich, aber einfacher

## Shared Variable ◁ Modelling

---

- Synchronisation is about exchanging information, i.e. sharing memory
- Very small information → Timing dominated by memory access latency
- Memory access is cached → We have to model caching
- We identify a shared variable with the cache line it resides on
- MSI protocol + atomic operations

14

- kein/wenig “Rechnen”

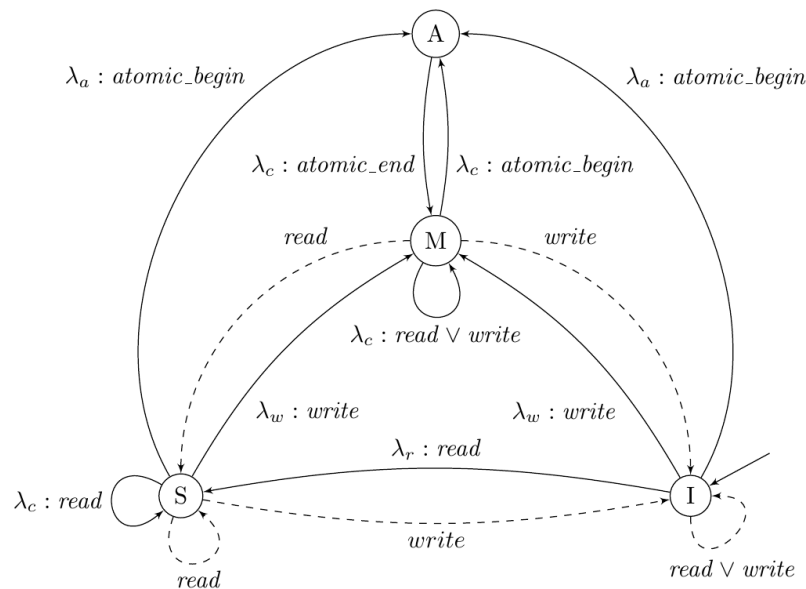
- Cache line Erkl  rung

- 64 byte Einheit von Speicher die am Stueck eingelagert und verdraengt werden.

- Adressen genau auf 64byte ausgerichtet

- Eine Kopie pro Core/Cache. Synchronisation zwischen den Kopien noetig damit keine Veralteten Daten benutzt werden. Beispiel: Jemand schreibt und du benutzt ein altes Datum.

# Shared Variable $\triangleleft$ Modelling




15

- CTMC Modul
- Ein Model setzt sich aus mehreren CTMC Modulen, die ueberlappt ausgefuehrt werden. Synchronisation zwischen Modulen ueber action labels.
- Eine cache line Kopie pro thread und pro Variable
- Synchronisation zwischen allen Modulen der gleichen Variable ueber action labels
- $\lambda_c = 1$  cycle
- $\lambda_r = 50$  cycles
- $\lambda_w = 100$  cycles
- gestrichelte Linien bekommen ihre Rate von der ausloesenden Aktion auf der ausloesenden cache line Kopie



## Central Counter Bar. $\triangleleft$ Modelling

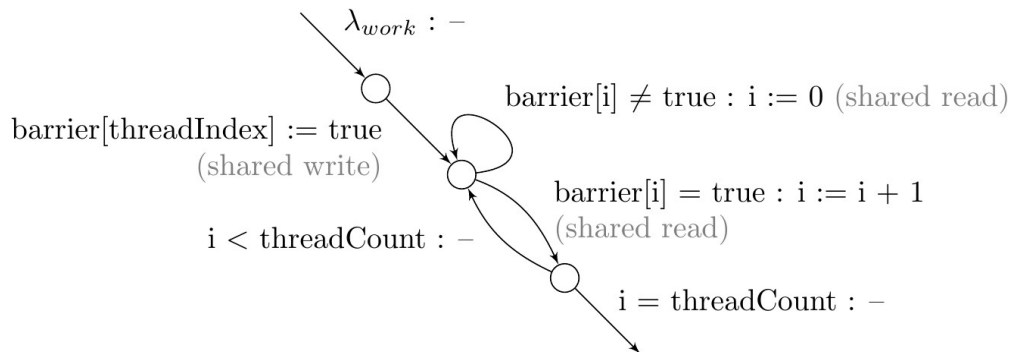


16

- Ein Protokoll-Modul pro thread
- Ein cache line Kopie Modul pro Thread, weil nur eine Variable
- Rate von shared memoryOperationen sind abhaengig vom cache line Kopie Zustand der jeweiligen Variable
- work ist Verteilung der Ankunftszeiten der threads an der Barriere
- Ungleichgewicht da Arbeit unterschiedlich schnell fertig gestellt wird

## B1 Barrier $\triangleleft$ Modelling

---



17

- Ein cache line Kopie Modul fuer jedes Array Element
  - n Protokol Module
  - $n \cdot n$  cache line Kopie Module

## Functional $\triangleleft$ Analysis

---

- “A thread may only exit the barrier if all threads have entered it”

$$\Box(one\_left \implies all\_entered)$$

- “If all threads entered the barrier, each one leaves it in a finite amount of time”

$$\Box(all\_entered \implies \Diamond all\_left)$$

## Functional $\triangleleft$ Analysis

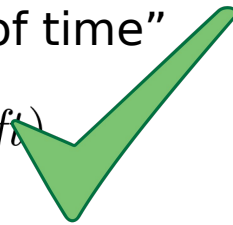
---

- “A thread may only exit the barrier if all threads have entered it”

$$\Box(one\_left \implies all\_entered)$$

- “If all threads entered the barrier, each one leaves it in a finite amount of time”

$$\Box(all\_entered \implies \Diamond all\_left)$$



## Quantitative ◁ Analysis

---

- Interesting values
  - time in cycles (2.5GHz clock speed)
  - energy consumption in joule / watt
  - implemented using CTMC time and rewards
- Parameters to variate
  - Number of threads
  - Rate of the initial (work) distribution

20

- kurz rewards erklären
  - transitions system bekommt 2 neue Funktionen
  - state, trans reward ...

## Quantitative $\triangleleft$ Analysis

---

- Points in time when to measure
  - (A) First thread entered the barrier
  - (B) Last thread entered
  - (C) First thread left
  - (D) Last thread left
  - (W) Last thread finished writing
- We take the reachability reward wrt the proper reward functions at these points in time

21

- (A), (B) bezogen auf Arbeitsperiode ueberwinden

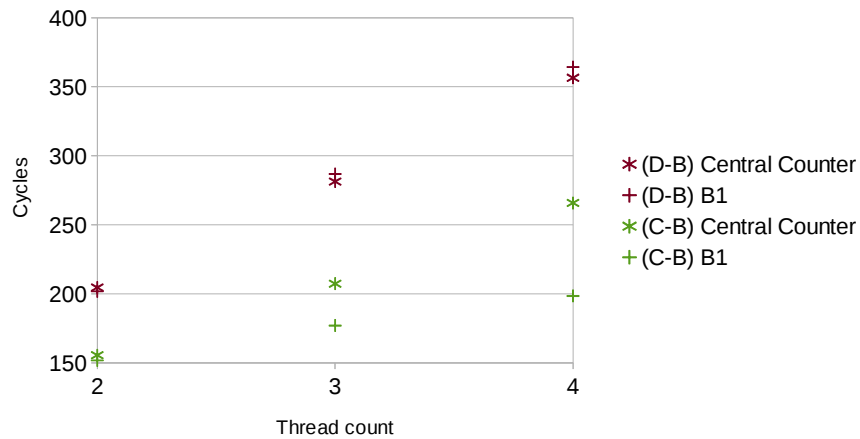
- Average reward

- “Proper” = reward functions die die richtigen Werte repraesentieren fuer Zeit und Energieverbrauch

11 Watt baseline. 2 nJ for fast shared mem ops.  
200 nJ for slow shared mem ops

## Quantitative ◁ Analysis

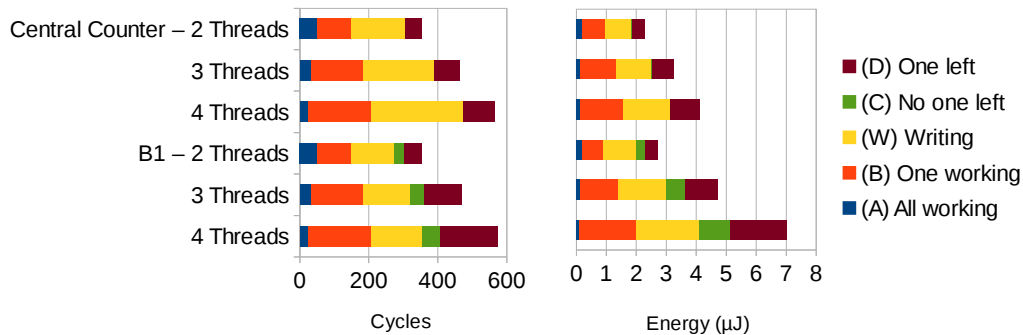
- small work period (100 cycles)



22

- C-B letzter der betritt bis zum ersten der verlaesst
- D-B letzter der betritt bis zum letzten der verlaesst
- 150 cycles minimum (2 shared writes)

## Quantitative Analysis

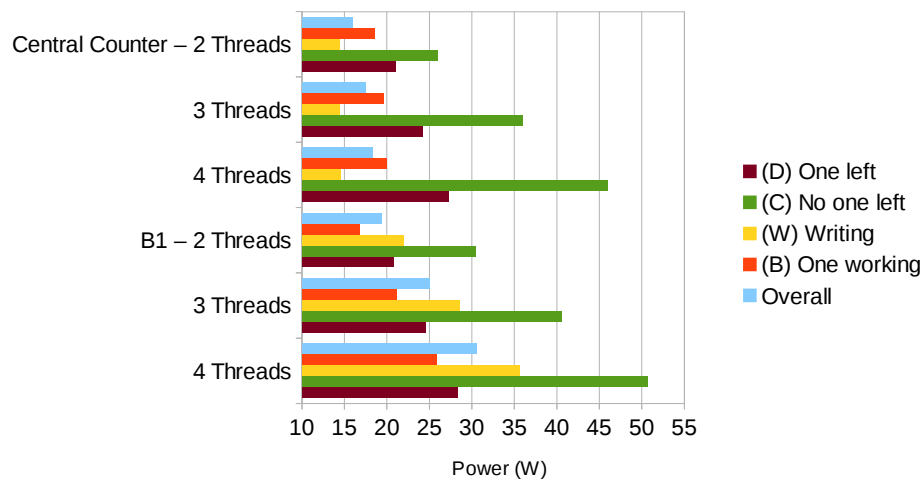


23

- A B W C D nochmal erklären (andere Wortwahl in der Legende)
- Energie ähnlich wie Timing
- atomic op blockieren spart Energie
- A, B nur abhängig von der thread Anzahl
- kein Gruen fuer CC, weil letzter ankommer verlaesst die barrier (mit hoher wahrschl.) sofort
- Mit steigender thread-Anzahl steigt die Laufzeit aller Operationen wegen der CTMC Semantik zu gleichzeitig aktivierten Transitionen
  - semantik erklären
  - zB CC Phase (D) 50, 75, 92 cycles
  - Erwartung waere 50 fuer alle
  - zB Phase A =  $100 / n$ , B =  $\sum_{i=1}^n 100/i$
  - deswegen dauert (W) 125, 138, 146 Takte statt 100 ueberall wie erwartet
- Lesen / Schreiben ... Vermutung reiterieren.




## Quantitative $\triangleleft$ Analysis



24

- (A) has the 11 watts baseline
- overall CC 16-18.3 watts, B1 19.4-30 watts
- (B) wie erwartet ... busy waiting / concurrent write / atomic ops
- (W) atomic ops + busy wait vs concurrent write + busy wait
- (C) sehr kurz / viele shared reads (verbrauchen mehr energie)
- (D) aehnlich, weil bei beiden nur lesen

## Quantitative $\triangleleft$ Analysis

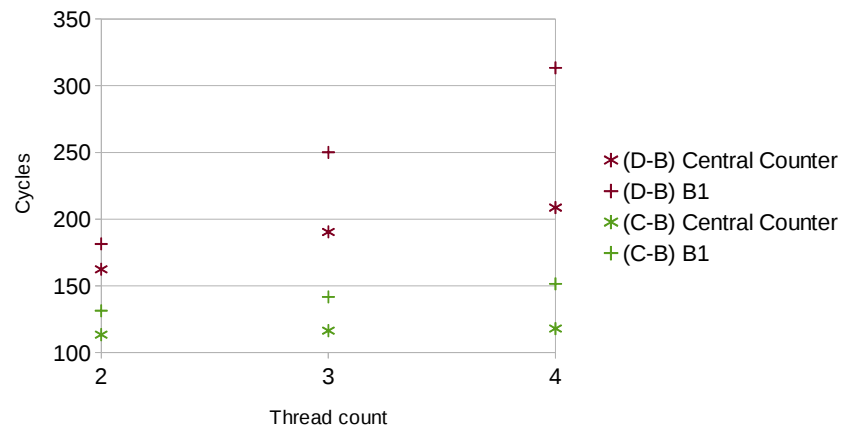


25

- $W = (W)$  = first enters to last wrote
- $R = (D-W)$  = after  $W$  until finish
- Erwartung CC schreiben wird groesser ... passiert nicht wegen CTMC Semantik ((D) wird groesser mit steigendem core count wie vorher erklart)
- B1 wie erwartet

# Quantitative $\triangleleft$ Analysis

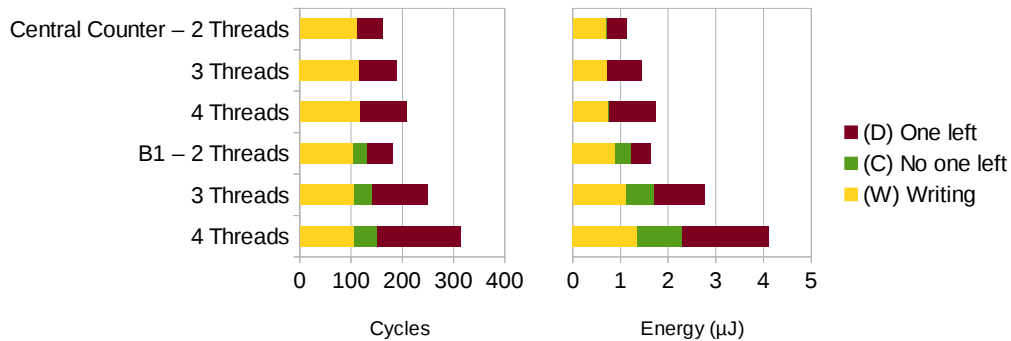
- large work period (1000 cycles)



26

- keine/wenige ueberlappende schreib operationen
- + langsames lesen -> alles langsamer

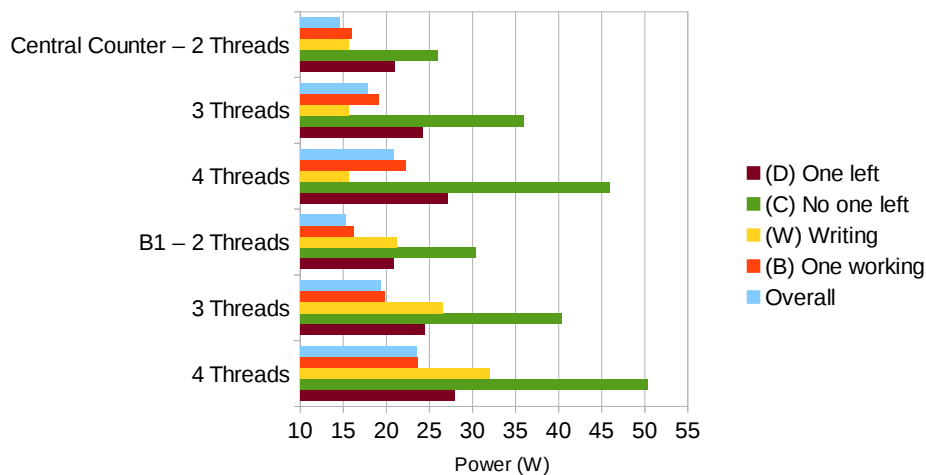
# Quantitative $\triangleleft$ Analysis



27

- (W) CC = 1 atomic op
- (W) B1 = 1 write
- (C-W) and (D-C) gleich wie bei 100 Takten, weil diese unabhaengig von der Eintrittszeit ist.
- (Sobald alle drinnen sind faengt ist W zu Ende und C faengt an)
- Energie (W) B1 steigt (Zeit nicht) weil mehr threads -> mehr busy waiting

## Quantitative $\triangleleft$ Analysis



28

- overall ausgeglichener weil grosse Arbeitsperiode mit 11 watt
  - CC = 14,6 - 20.9
  - B1 = 15.3 - 23.5
- B, C, D sehr aehnlich zwischen beiden
- W anders weil lesen und schreiben blockiert ist waehrend atomic ops -> nicht so bei B1
- je mehr Arbeit desto weiter naehern sich der Energieverbrauch zwischen den beiden Protokollen an

# Conclusion

---

- Introduced innovative barrier protocols
  - + No atomic operations or locks required
  - + Competitive performance
- Principles of pW/CS locks apt to improve synchronisation performance
- Quantitative model checking enables exhaustive, fine-grained analysis beyond the capability of tests and benchmarks

## Future Work

---

- Analyse protocols using measurement
- Invent more protocols
  - Variations of existing
  - Remote write-based
- Extend model checking
  - More processes/threads
  - More detail
    - Cache protocols, cache hierarchies
    - Limited bandwidth and other influences

# Sources

---

- [1] Probabilistic write copy select, Paper,  
In 13th Real-Time Linux Workshop, pages 195–206, Oct. 2011
- [2] A probabilistic quantitative analysis of  
probabilistic-write/copy-select, Paper,  
In NASA Formal Methods, pages 307–321. Springer, 2013.
- [3] Evaluation of publicly available Barrier-Algorithms and  
Improvement of the Barrier-Operation for large-scale  
Cluster-Systems with special Attention on InfiniBand Networks  
<http://hlor.inf.ethz.ch/publications/index.php?pub=12>
- [4] PRISM, Website, 13-03-019  
<http://www.prismmodelchecker.org>
- [5] SPIN, Website, 13-01-08  
<http://spinroot.com>



## Not Covered in the Presentation

- Survey of means to implement barriers
- Barrier building blocks
- Dissemination Barrier, MGB and B2 Barrier
- more in-depth analysis of the protocols

# Thank you!

Slides and report are available at

<http://automaton2000.com/barrier-slides.pdf>

<http://automaton2000.com/barrier-minor-thesis.pdf>